

Designing *Award-Winning* UI with Opus 4.8

How to aim Claude so the output is professional, distinctive, jury-grade work — the leverage points, the prompt anatomy, the review pipeline, and a copy-paste prompt library. Grounded in Anthropic's own Opus 4.8 guidance and the Awwwards scoring rubric.

00 The one page that matters

If you read nothing else.

Opus 4.8 is an unusually strong *executor* of visual design and a weak *originator* of taste. Your job is not to describe a vibe and hope — it is to supply the taste (a concrete reference + a non-negotiable spec), set the effort high, and then run a tight see-critique-fix loop until it clears a named bar. Award-winning is not a prompt; it's a pipeline.

LEVER	THE MOVE
1 · Effort	Use Opus at high/xhigh for any real design work. This is the single biggest quality knob in 4.8 — bigger than prompt wording.
2 · Reference	Give an artifact , not adjectives. A screenshot, an IDE theme, a named cultural aesthetic, a competitor frame to borrow composition from. "Make it clean" gets you slop.
3 · Spec	4.8 is literal . Pin exact hexes, type, radius, spacing, scope. Vague negations ("not generic") just shift it to a different generic.
4 · Diverge first	Make it propose 3–4 distinct directions before building . Breaks its default house style and hands you the steering wheel.
5 · The loop	I can see my own output (browser screenshots). Always make me screenshot, critique against the rubric, and fix — never accept first render.
6 · Adversary	End every slice with a fresh-eye critique (a design-review pass, Codex, or a cold sub-agent). The maker can't see the slop.

THE SHAPE OF A GOOD REQUEST

One upfront, fully-specified turn beats a vague brief dribbled across ten. 4.8 rewards autonomy: state task + audience + feeling + reference + constraints + the done-when bar in the first message, then let it run and verify itself.

• Contents

1. 01 The mental model — what 4.8 is, and isn't, good at
2. 02 The Opus 4.8 operating manual (the model-specific facts)
3. 03 The five highest-leverage moves
4. 04 The prompt anatomy — a reusable template
5. 05 The pipeline — phase gates mapped to your tools
6. 06 What "award-winning" actually means (the rubric)

- 7. 07 Anti-patterns & AI-slop tells
- 8. 08 Copy-paste prompt library
- 9. A Applied: the graph-paper UI track
- 10. B Sources

01 The mental model

Aim the tool at what it's actually good at.

The most common way to waste Opus on design is to treat it like a creative director — "design me something beautiful." It will, dutifully, regress to the mean of everything it has seen: competent, plausible, forgettable. The model is a **world-class implementer with strong-but-conventional instincts**. Give it taste and it executes at a level most teams can't; leave taste unspecified and it fills the vacuum with the average of the internet.

OPUS 4.8 IS EXCELLENT AT	OPUS 4.8 IS WEAK AT (YOU MUST SUPPLY)
Translating a concrete reference into clean, consistent code	Originating a distinctive point of view from a vague brief
Holding a design system across many components	Knowing when "correct" is still boring
Spacing/type scales, alignment, responsive logic, accessibility	Taste calls — the 5% that separates good from award-grade
Critiquing a rendered screenshot against a rubric	Trusting its own first answer is the best one
Generating many variants fast	Picking the brave variant over the safe one

Your leverage is in three places: the reference you feed in, the bar you hold it to, and the number of critique loops you force. Everything in this doc is a way to maximize those three.

The corollary: **convergence is a feature you turn on at the end, not the start**. Early, push for divergence and bravery (many directions, unexpected references). Late, push for consistency and polish (one system, ruthlessly applied). Mixing these — asking for "creative" and "consistent" in the same breath — is why output feels both random and bland.

02 The Opus 4.8 operating manual

Model-specific behavior, straight from Anthropic's 4.8 guidance — these change how you prompt.

Effort is the master knob

4.8 respects the effort/thinking setting *strictly*, more than any prior Opus. At low/medium it scopes narrowly and can under-think; design work wants **high or xhigh**. If you see shallow output, the first move is "raise effort," not "rewrite the prompt."

IN YOUR ENVIRONMENT

In Claude Code you're on Opus 4.8 already (your Max 20× home base). `/fast` gives *faster output at the same intelligence* — it does not downgrade the model — so leave it on for iteration without fear. For genuinely hard design reasoning (a from-scratch system, a gnarly layout), give it room: don't interrupt, let it think.

It is literal now — specify scope

4.8 does not silently generalize. "Tighten the spacing" fixes the one spot it judged worst. If you mean everywhere, say "*apply to every section, every breakpoint.*" This literalism is a gift for pipelines (predictable) and a trap for lazy briefs.

Its default "house style" — know it, then override it

Left to itself, 4.8 reaches for a recognizable look: **warm cream backgrounds** (~ `#F4F1EA`), **serif display** (Georgia / Fraunces / Playfair), **italic word-accents**, **terracotta/amber accent**. Beautiful for editorial, hospitality, portfolios — wrong for dashboards, dev tools, fintech, anything that should feel engineered.

LUCKY BREAK FOR THIS PROJECT

That default is *almost exactly* the graph-paper aesthetic you already want — warm paper, ink, hand-serif. For Immortal Engine you're working **with** the grain, which is rare. But the default still drifts to *generic* warm. You must pin the *specific* paper (`#e7ecdd`), the quadrille grid, the ballpoint-ink palette — or it will quietly become Fraunces-on-cream like everything else it makes.

Critically: **generic negations don't work**. "Don't use cream / make it minimal" just swaps one fixed palette for another. Two things that *do* work: (1) give a concrete alternative spec, (2) make it propose options first (next section).

One rich turn > many thin turns

4.8 is built for autonomy. A single, fully-specified opening message produces better and cheaper results than the same information revealed across many back-and-forths. Front-load intent, constraints, and the done-when bar.

Reviews: it follows your bar faithfully

4.8 is better at finding issues but also more obedient to "only report high-severity." For a design critique you want *coverage*: tell it to report everything, including low-confidence and nit-level, and rank by severity — let a second pass filter. (Used in §08's audit prompt.)

03 The five highest-leverage moves

Where effort actually converts to award-grade output.

1 — Reference-driven, not adjective-driven

Adjectives ("modern, clean, premium") are noise; every brand claims them. Replace each with a *pointer*:

- **An artifact** — a screenshot/URL/image of the exact thing you admire (or of your own current state, for "fix this").
- **A theme transfer** — "borrow the palette of the Nord IDE theme," "the typography of a Penguin Classics cover," "the grid of Swiss railway signage."
- **A cultural anchor** — a named aesthetic (brutalist, solarpunk, 1970s technical manual, quadrille engineering pad) with 3–4 concrete attributes.
- **A composition donor** — "borrow the layout rhythm of [Dribbble/site], none of its colors." Steals craft, keeps it yours.

2 — Diverge before you build

The most reliable single technique for beating the default and getting bravery: force option-generation first.

PASTE THIS BEFORE ANY "BUILD IT" REQUEST

Before writing code, propose 4 distinct visual directions for this brief.
For each: background hex / accent hex / typeface / one-line rationale +
the one risk that makes it interesting. Make them genuinely different
(not 4 shades of the same idea). I'll pick one, then you build only that.

3 — Direct one dimension at a time

Quality lives in dimensions, and they're easier to nail individually. When polishing, name the axis: **typography, color, motion, background, spacing, hierarchy**.

DIMENSION WHAT "GOOD" LOOKS LIKE (AND WHAT TO ASK FOR)

Typography	One distinctive face used decisively. Avoid Inter/Roboto/Arial/system . High-contrast pairing (display + mono, serif + geometric sans). Weight extremes (200 vs 800, not 400 vs 600). Size jumps 3×+, not 1.5×.
Color	A <i>dominant</i> color + sharp accent beats a timid even spread. Lock as CSS variables. Steal from IDE themes. Never purple-gradient-on-white.
Motion	<i>One</i> well-orchestrated page-load with staggered reveals (animation-delay) delights more than scattered micro-interactions. CSS-first.
Background	Atmosphere & depth, not a flat fill — layered gradients, geometric pattern, a structural texture (e.g. your quadrille grid). Make it earn its place.
Spacing	A real scale (4/8px base), generous whitespace, things that align to a grid. This is 50% of "looks professional."

4 — Close the visual loop (you already have the rig)

The thing that makes me reliably good at UI is that **I can see what I made**. Your repo already has the preview/browser harness and the playtest screenshot discipline. Use it as a hard rule: *edit* → *reload* → *screenshot* → *critique* → *fix*, and never let me declare done from code alone. "Verify with a screenshot before calling it done" should be in every UI instruction.

5 — Always bring an adversary

The maker is blind to its own slop. End each slice with a cold critique — a different lens, ideally a different context window: your `design-review` skill, a Codex review, or a fresh sub-agent told to be a harsh Awwwards juror. Three independent eyes (build, design-critique, code-critique) is the cheapest quality multiplier you have.

04 The prompt anatomy

A reusable skeleton. Fill all seven slots; missing slots are where slop enters.

SLOT	WHAT GOES HERE
1 Role+altitude	"You're a senior product designer. Taste bar: Awwwards SOTD." Sets the ceiling.
2 What & who	The exact surface + the user + the moment they're in. ("The in-play screen; a solo player mid-session, reading.")
3 Feeling	3–4 adjectives <i>plus the reference that proves them</i> .
4 Reference	The artifact / theme / cultural anchor (move #1). Non-optional.
5 Constraints	The hard spec: exact hexes, type, radius, spacing scale, what NOT to touch, the invariants.
6 Done-when	The measurable bar: "reads as one graph-paper sheet on desktop+mobile; zero console errors; verified by screenshot."
7 Output contract	Diverge-first? One file? Show before/after? Self-critique before finishing?

WORKED EXAMPLE – A SLICE OF YOUR REAL WORK

ROLE You're a senior product designer; bar = Awwwards Site of the Day craft.

TASK Redesign the Immortal Engine in-play screen (public/v1.html, v1.js, styles.css). User: a solo player mid-session, reading and typing to a DM.

FEEL A real graph-paper engineering pad someone is playing on – handmade, inky, lived-in. Reference: the canvas map already does this exactly (LocalMap.js PAPER #e7ecdd, teal quadrille, blue ballpoint). Make the WHOLE sheet that, with the map as one inked region. Anti-reference: a dark fantasy "app with panels" dashboard.

SPEC Paper #e7ecdd; grid 24px minor / 120px major rgba(92,134,120,.17/.34); ink #222c48; pen accent #2a4a78; red felt-tip #b3331f. Hand-letter section labels. Ruled ink boxes, no glossy cards, no shadows, no Diablo orbs, no system quest/rumor tracker. Type stays readable.

GUARD Presentation only – do NOT touch engine/server/world state; worldHash unaffected; revertable by file. Preserve the DM-test + two-surface rule.

DONE The play screen reads as one continuous graph-paper artifact, desktop and mobile; zero console errors. Verify with screenshots before done.

OUT First propose 3 texture directions (clean-ruled / taped-scrap / field-notebook) with a one-line risk each; I pick one; then build.

Note how slots 4–6 do the heavy lifting. The reference makes it distinctive; the spec makes it precise (4.8 is literal); the done-when makes it self-verify. Strip those three and you'd get generic-warm slop.

05 The pipeline

Award-grade is a sequence of gates, not one shot. Mapped to skills you already have.

PHASE	GOAL & HOW TO INVOKE IT
0 · System	Define the design language once. Run <code>/design-consultation</code> → a <code>DESIGN.md</code> (palette, type scale, spacing, motion, voice) that becomes the source of truth every later prompt cites. Do this before pixels. For an existing look, <code>/plan-design-review</code> infers the system instead.
1 · Diverge	Explore directions, don't commit. <code>/design-shotgun</code> generates multiple variants on a board to compare; or use the diverge-first prompt (§03 #2). Pick the brave one.
2 · Build	Implement the chosen direction as a small, revertable slice. Direct edits, or <code>/design-html</code> for production-quality HTML/CSS from an approved mockup. One coherent change per slice.
3 · Verify	See it run. Preview server → reload → screenshot (desktop + mobile) → check console. This is non-negotiable and already your playtest protocol. No "done" without a screenshot.
4 · Critique	Adversarial review. <code>/design-review</code> (designer's-eye visual QA, fixes in-source with before/after) + a <code>/codex</code> review for code health + optionally a cold sub-agent as a harsh juror. Coverage over politeness (§02).
5 · Iterate	Loop 3↔4 against the §06 rubric until it clears the bar. Each pass: one dimension (§03 #3). Commit the slice; move to the next.

SHORTCUT

`/autoplan` runs the CEO / design / eng / DX reviews in sequence and only surfaces the taste calls to you at the end — good when you want the full gauntlet on a plan without answering 30 intermediate questions. Use it on the *plan*, before building, for anything ambitious.

Rule of thumb: time-box building, never time-box critiquing. The award lives in loops 4–5, the part everyone skips.

06 What "award-winning" actually means

Borrow the jury's rubric so I can hold myself to it. Awwwards' real weights:

CRITERION	WEIGHT	WHAT IT REWARDS
Design	40%	Visual craft: color, typography, imagery, layout, consistency, the details.
Usability	30%	Navigation, clarity, load speed , responsiveness across devices.
Creativity	20%	A fresh idea — original concept, inventive use of the medium.
Content	10%	Quality & relevance of the actual words and material.

Sites go to ≥ 18 jurors; the 3 scores furthest from the average are dropped. Translation: **polish is table stakes (70% is design+usability), but the win is one brave, coherent idea executed without a single sloppy detail.** Consensus competence scores ~ 7.5 ; a singular vision with flawless craft scores 9+.

THE CRAFT CHECKLIST (WHAT JURORS NOTICE, AND WHAT TO MAKE ME AUDIT)

- Type scale is intentional ($3\times +$ jumps), one distinctive face, real hierarchy
- Spacing on a strict $4/8$ grid; nothing "almost" aligned
- Dominant color + sharp accent; consistent CSS variables; no muddy middles
- Optical alignment, consistent corner radii, consistent icon weight
- One orchestrated entrance animation; purposeful motion, never decorative jitter
- Background has depth/texture, not a flat fill
- Empty/loading/error/long-text states all designed, not just the happy path
- Responsive: genuinely re-laid out at mobile, not just squished
- Accessibility: contrast $\geq 4.5:1$, focus states, hit targets $\geq 44\text{px}$, reduced-motion
- Performance: fast first paint, no layout shift, no jank
- Microcopy has voice; no lorem, no "Welcome to our app"
- The 5%: one detail that makes someone smile and screenshot it

07 Anti-patterns & AI-slop tells

If you see these in my output, the brief failed upstream. Name them and send it back.

PROMPT-SIDE MISTAKES (YOURS TO AVOID)

- **Adjectives without a reference.** "Sleek and modern" $\rightarrow$ average of the internet.
- **Generic negations.** "Don't make it generic / not too busy" $\rightarrow$ swaps one default for another.
- **Accepting the first render.** The first is the safe one. Always loop.
- **No done-when bar.** Without a finish line I stop at "plausible."
- **Thin multi-turn dribble.** Front-load the spec; 4.8 pays you back in autonomy.
- **Skipping the screenshot.** Code that compiles $\neq$ UI that looks right.

- **Letting me invent scope.** A landing page you didn't ask for, a new tracker panel — pin scope or 4.8's literalism cuts the other way and it builds what it imagined.

OUTPUT-SIDE TELLS (THE SLOP SIGNATURE)

- Inter / Roboto / Arial / system fonts; type with no contrast
- Purple (or any) gradient on white/dark; timid evenly-spread palette
- Everything centered; identical card grid; rounded-corner sea of boxes
- Drop-shadows doing the work color/hierarchy should do
- Emoji as iconography; "🌟 Welcome" energy
- Decorative motion with no orchestration; hover-jitter everywhere
- Flat solid background; no atmosphere or texture
- Dark-mode-by-default as a substitute for taste

THE TELL BEHIND THE TELLS

Slop is *contextless competence* — it would look the same for any product. The cure is always the same: more specific context, a real reference, and a higher bar. Distinctiveness comes from the brief, not from asking harder for "creativity."

08 Copy-paste prompt library

Eight ready prompts. Combine with the anatomy in §04.

A • THE STANDING AESTHETICS GUARDRAIL (PUT IN CLAUDE.MD)

ADD TO PROJECT INSTRUCTIONS, ONCE

<frontend_aesthetics>

NEVER use generic AI aesthetics: overused fonts (Inter, Roboto, Arial, system), cliched schemes (esp. purple gradients on white/dark), predictable layouts, cookie-cutter components. Use a distinctive face, a dominant color + sharp accent as CSS variables, atmospheric/textured backgrounds, and one orchestrated load animation. For any UI change: edit → reload → screenshot (desktop + mobile) → critique against the Awwwards rubric (design/usability/creativity/content) → fix. Never declare a visual change done without a screenshot. Presentation work must not touch engine/server/world state.

</frontend_aesthetics>

B • KICK OFF A DESIGN SYSTEM

/design-consultation

Build DESIGN.md for Immortal Engine's player UI. North star: the whole interface is one continuous graph-paper engineering pad (ref: the canvas map, #e7ecdd quadrille, blue ballpoint ink). Define palette, type scale, spacing scale, the "ruled-box" component language, motion, and DM voice. Cite the two-surface rule and DM-test as constraints.

C • DIVERGE FIRST

Before building, propose 4 genuinely distinct directions for <surface>. Each: bg hex / accent hex / typeface / texture / one-line rationale + the risk that makes it interesting. No two may share a palette. I pick one; build only that.

D • BUILD A SLICE

Implement direction <N> on <files> as ONE revertable slice. Hard spec: <hexes/type/radius/spacing>. Do not touch <invariants>. When done, reload the preview, screenshot desktop + 375px mobile, check console, and show me the screenshots. Don't call it done if anything looks off – fix first.

E • SINGLE-DIMENSION POLISH

Polish ONLY typography on <surface> (leave color/layout untouched): one distinctive face, high-contrast pairing, weight extremes, 3x size jumps, true hierarchy. State the choice before coding, then show a before/after screenshot.

F • THE HARSH-JUROR AUDIT

Act as an Awwwards juror reviewing <surface> for Site of the Day. Score design/usability/creativity/content (0–10 each, weighted 40/30/20/10) from a screenshot. Report EVERY issue incl. low-severity and uncertain ones, with severity + the specific fix. Coverage, not politeness – a later pass will filter. End with: the single change that would most raise the score.

G • SPIN OFF AN INDEPENDENT CRITIQUE

/design-review (then) /codex review

Two independent lenses on the rendered slice — visual craft, then code health. For a third, ask for a cold sub-agent told to be a brutal juror with no memory of building it.

H • ACCESSIBILITY & RESPONSIVE GATE

Audit <surface> for: contrast $\geq 4.5:1$ (list failures w/ hexes), visible focus states, hit targets $\geq 44\text{px}$, prefers-reduced-motion, and a genuine mobile re-layout (not squish). Fix each, re-screenshot at 375 / 768 / 1280px, and confirm.

A Applied: the graph-paper UI track

Turning this doc on the work already in motion.

You've shipped Slice 1 (the substrate) — the whole play screen now reads as one graph-paper sheet. Here's how to run the rest at award grade rather than "fine."

DO THESE NOW

1. Write **DESIGN.md** (prompt B). Right now the spec lives in **docs/GRAPH_PAPER_UI.md** as prose; promote the palette/type/spacing/"ruled-box" rules into a real design system every future slice cites. This is the highest-ROI next step — it stops drift.
2. **Reconcile the one visible seam:** the canvas map's paper tint vs the page grid. Either match them exactly, or deliberately frame the map as a sketch boxed on the sheet. Decide on purpose — jurors notice the near-miss.
3. **Diverge on texture** (prompt C) before Slice 2: clean-ruled vs taped-scrap vs field-notebook-with-fold. Pick the brave one; don't default into "clean."
4. **Add the standing guardrail** (prompt A) to **CLAUDE.md** so every UI turn self-verifies by screenshot — you already do this in playtest protocol; make it law for visual work too.

WHERE THE AWARD ACTUALLY IS, FOR THIS PRODUCT

The brave idea is already yours and it's strong: *the interface is a physical artifact the player writes on, and it only knows what the player has witnessed*. That's the 20% Creativity score most products never get. To convert it, the execution has to be flawless in the 40% Design / 30% Usability:

- **Make it feel handmade, not CSS** (Slice 2): slight stroke wobble, a fold crease, one taped scrap — texture that says "a person drew this."
- **Hand-letter consistently** — labels, headings, marginalia in one voice; typeset only where reading demands it (narration).
- **Combat as marginal marks** (Slice 4) is the chance for the screenshot-worthy 5%: pencil rings, struck-through enemies, a d20 result jotted in the margin that erases after the fight.
- **Kill the trackers** (dead goal CSS, the rumor board) — a clean sheet with nothing systemic on it is itself the statement.
- **Usability guard:** readability on the busy grid is your risk. Hold the contrast/responsive gate (prompt H) every slice.

THE HONEST CONSTRAINT

You're inside a 2,900-line vanilla **v1.js** with live game state. "Award-winning" here means *the rendered play screen could win a craft award*, achieved through presentation-only slices that never risk the engine — exactly the discipline you set in **GRAPH_PAPER_UI.md**. Don't let a redesign ambition pull edits into the engine; the whole point is that taste and determinism live in different layers.

B Sources

Primary, current (fetched 2026-06-26).

- Anthropic — *Prompting Claude Opus 4.8* (effort, literalism, default house style, autonomy, review behavior). platform.claude.com/docs/.../prompting-claude-opus-4-8
- Anthropic Cookbook — *Prompting for frontend aesthetics* (the dimension strategies, the distilled aesthetics prompt, anti-patterns). platform.claude.com/cookbook/coding-prompting-for-frontend-aesthetics
- Anthropic — *Claude prompting best practices*. platform.claude.com/docs/.../claude-prompting-best-practices
- Awwwards — *Evaluation System* (Design 40 / Usability 30 / Creativity 20 / Content 10; ≥18 jurors, outlier-trim). awwwards.com/about-evaluation
- Design Systems Collective — *From Prompt to Production: a designer's workflow with Claude Design + Claude Code*.
- Community references — *awesome-claude-design* (DESIGN.md prompt families); baoyu-design (local design skill for Opus 4.8).

Internal working document. The prompt blocks are starting points — the discipline in §02–§05 (effort high, reference in, bar named, loop forced, adversary brought) is what actually moves output from competent to award-grade. Pair with [docs/GRAPH_PAPER_UI.md](#) for the project-specific spec.